Interpreting Random Forest Models Using Feature Importance: A Case Study with the Titanic Dataset

GitHub: https://github.com/Manasa-2409/machine-learning-/tree/main

Student name: Manasa

Student ID: 24158060

Introduction

There is an increasing use of machine learning models in pattern finding and predictions. The value to understand the model's operation is as important as achieving a high predictive accuracy. Meaningful interpretation of machine learning models is essential for stakeholders in diverse fields like healthcare and finance to build trust in the models' output and gain knowledge from their data.

Random Forest is an ensemble learning algorithm that is widely used today owing to their powerful prediction capability and overfitting resistance. The algorithm generates various decision trees and combines their conclusions to form a final class. Due to their many trees, Random Forest models may seem too complex to interpret at first glance. Despite this disadvantage, Random Forest could compute the importance of features. This feature importance gives useful insight as to which features affect the predicted outcome the most.

By identifying the relative contribution of each variable, feature importance makes a complex ensemble model more interpretable and understandable. This tutorial demonstrates the use of feature importance for interpreting a Random Forest model. to do this we will use the Titanic dataset.

Through this tutorial, we will explore the interpretation of Random Forest models using feature importance techniques. In particular, the tutorial examines which passenger characteristics influence survival predictions the most and compares the unusual ways to measure feature importance.

Random Forest and Model Interpretability

Unlike a decision tree, a random forest is a combination of various trees. Which combine the predictions together. Each tree is trained on a randomly selected bootstrap sample of the training data. Also, during the construction of trees, split-of-features is considered only at a few features causing further randomness.

Through these mechanisms, various decision trees are formed that represent different patterns within the data. Aggregating the predictions of all the trees in the forest leads to the final prediction. In classification tasks, majority voting is typically used. Using this methodology, the variance is reduced and prediction performance is better than a decision tree.

While ensemble models are powerful, they can be criticized for being less interpretable than simpler models. Random Forest manages this issue, while providing feature importance estimates. Feature importance measures tell you how much impact each input variable has on the prediction performance of the model. Taking these scores into account, the analyst can figure out the most key features.

Dataset Description
The Titanic dataset contains information about passengers who travelled aboard the RMS Titanic, which sank in April 1912 after colliding with an iceberg. The dataset is commonly used for demonstrating classification algorithms because it contains a mix of demographic and socioeconomic variables that influence survival outcomes.

Key variables included in this tutorial are:

- Passenger class (pclass) – indicates socioeconomic status.
- Sex – gender of the passenger
- Age – age in years
- Siblings/Spouses aboard (sibsp) – number of relatives travelling together.
- Parents/Children aboard (parch) – number of family members on board.
- Fare – ticket price paid by the passenger.
- Embarked – port where the passenger boarded the ship.

The target variable is survived which shows whether the passenger survived or not. Consequently, the modelling problem may be framed as a binary classification problem in which the model predicts survival based on certain characteristics of a passenger.

Data Preprocessing was required before Model training. We will replace the missing values of the age variable with the median age and the missing values in the embarkation with the mode. The categorical variables sex and embarkation port were transformed into numerical using one hot encoding for them to be used by the machine learning algorithm.

Exploratory Data Analysis
Understanding the dataset is a major step before building a predictive model. The first step in the analysis involved examining the distribution of the target variable.

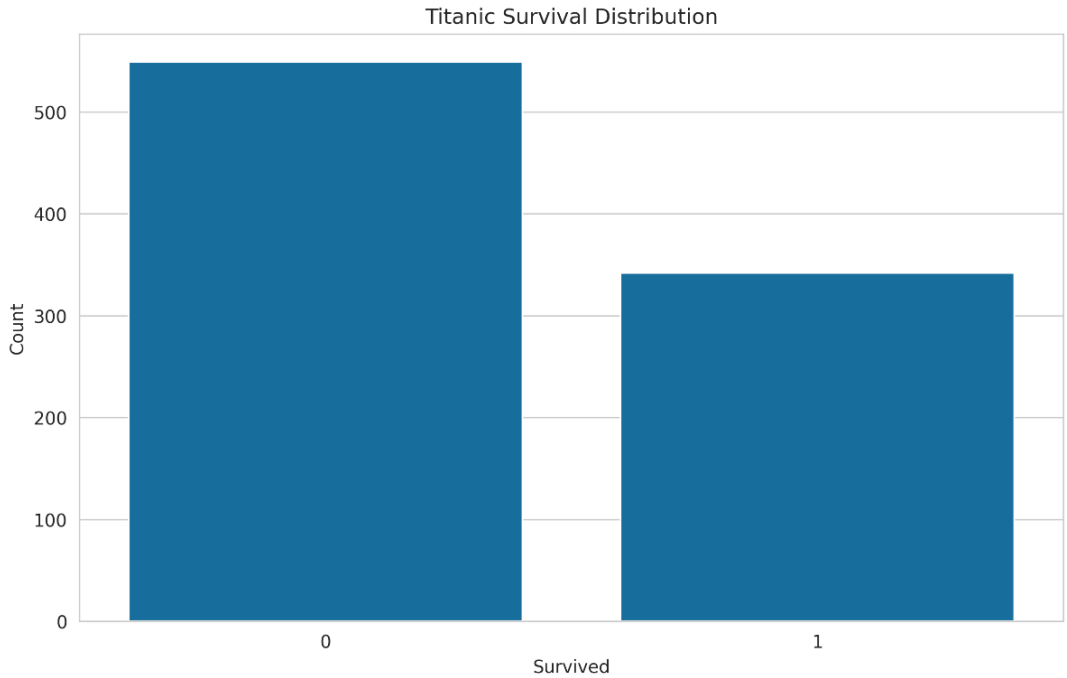

The distribution indicates that there are more non-survivors than survivors in this case. It is common for many real-world classification problems to exhibit this imbalance, which highlights the need for careful consideration of evaluation metrics.

Analysis of Titanic catastrophe as per historians points out that a passenger's gender & class very much affected survival possibility. Women and children were often evacuated before other passengers. Those in higher classes had better access to lifeboats. The historical background discussed above is a good context to keep in mind while interpreting the Random Forest results.

Model Training

After the dataset is pre-processed, Random Forest classifier is used to predict passenger survival. The model was created with 200 decision trees for stable predictions at a reasonable cost.

The Random Forest employs a collection of decision trees to analyse data and arrive at predictions. These trees predict the final classification by grouping their predictions. A combination of learning algorithms helps reduce overfitting and make better predictions.

After training of model, it was assessed on a separate test dataset to evaluate predictive performance. Accuracy was primarily used to assess how well the model predictions matched respect to survival.

Random Forest Feature Importance

One of the advantages of Random Forest models is the ability to estimate feature importance. In the default implementation, feature importance is calculated using the mean decrease in impurity. During the training process, decision trees repeatedly split the data using unique features to reduce classification error. Features that produce larger reductions in impurity across many trees receive higher importance scores.

By aggregating the importance values across all trees, the Random Forest model produces an overall ranking of features based on their contribution to prediction accuracy.

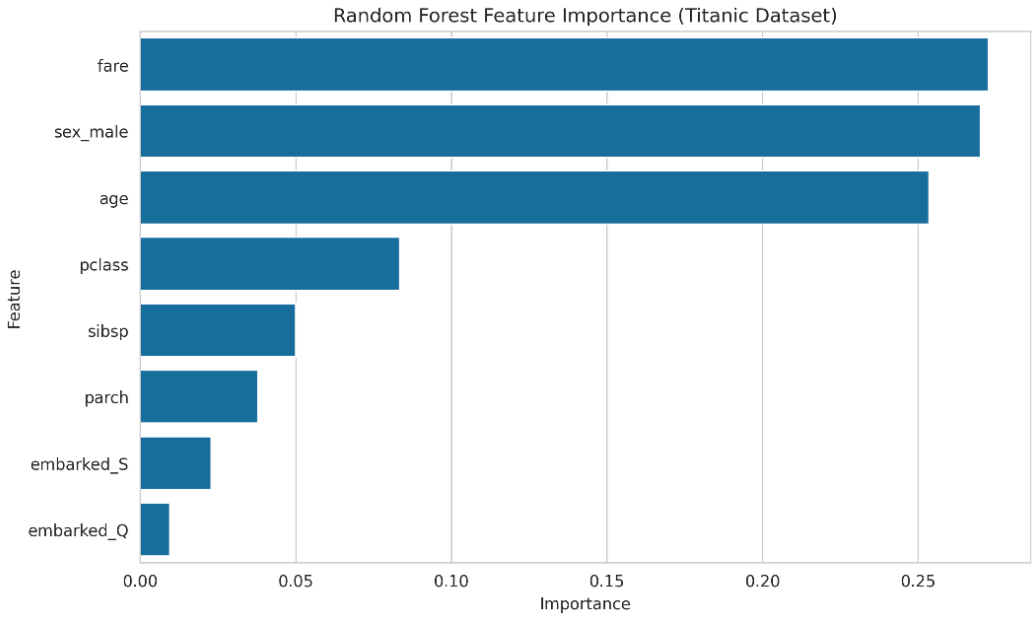

The results indicate that certain variables are more influential in predicting survival. Specifically, the most influential variables are gender and passenger class. The evidence reveals that women and people of higher classes had a better chance of survival.

Permutation Feature Importance

While the feature importance available in scikit-learn is helpful, they are not always the best estimates. It depends on whether the features are on the same scale and correlated. To overcome these limitations, we use one more method called permutation importance.

Permutation importance assesses the significance of a feature by analysing how the model's prediction accuracy decreases due to random shuffling of the feature's values. If the model’s accuracy drops considerably because of shuffling a feature, the feature is important.

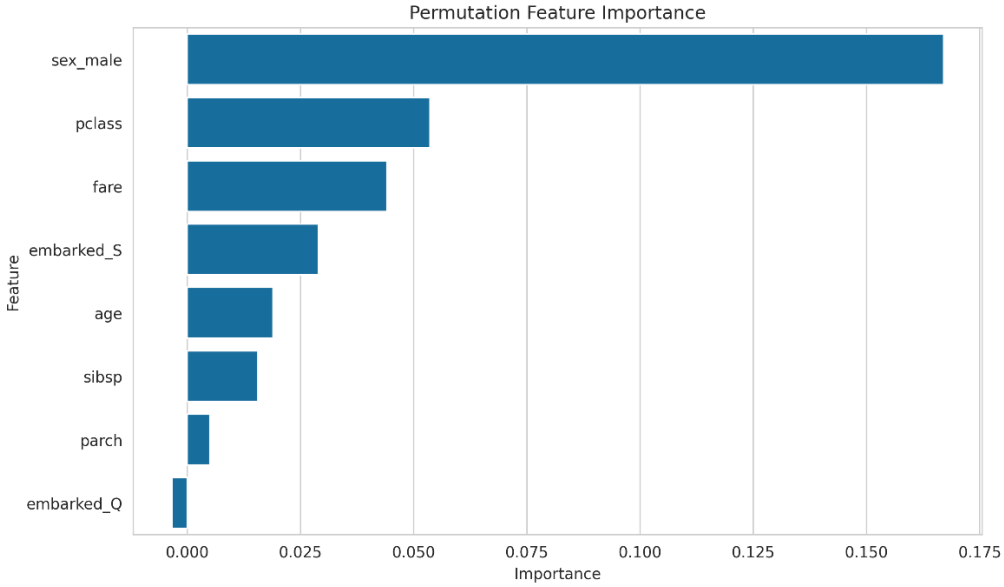

The results confirm that gender and passenger class remain among the most influential predictors of survival. This consistency across multiple importance methods strengthens confidence in the interpretation of the model.

Comparison of Feature Importance Methods

To further explore interpretability, the feature importance values obtained from the Random Forest model were compared with the permutation importance scores.

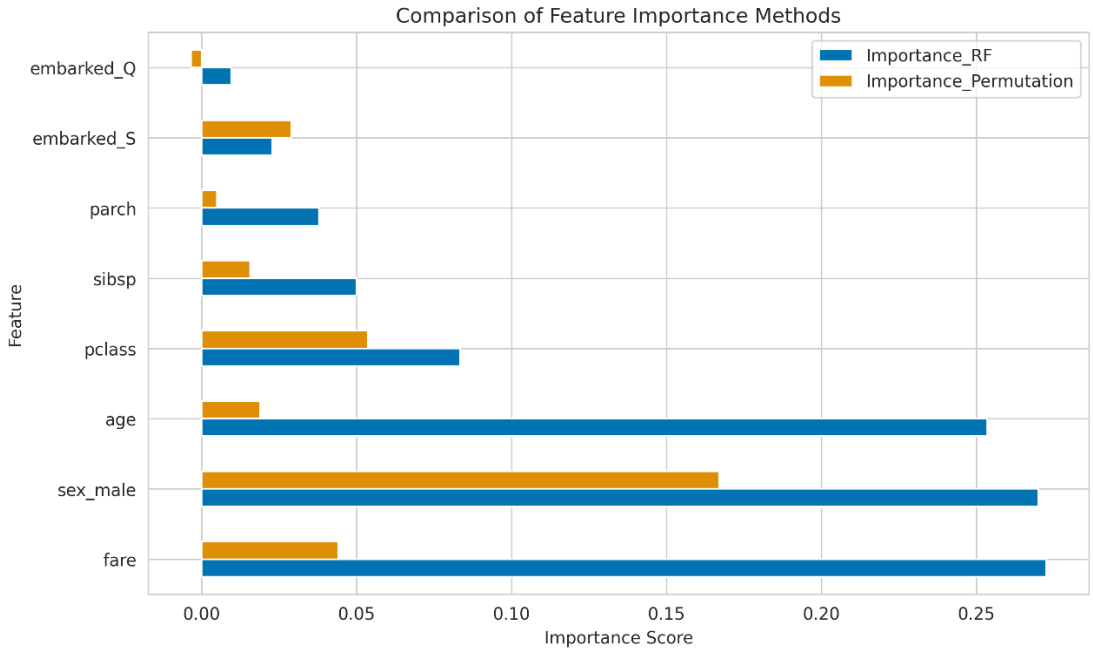

While both techniques reveal similar crucial features, slight differences can be seen in the ranking of variables. The techniques measure feature importance differently so there are such differences. The built-in Random Forest method focuses on the decrease in node impurity, whereas permutation importance indicates the effect of each feature on the accuracy of predictions.

By comparing these methods, we get to know the model in a deeper sense. Further having multiple interpretation techniques is important while dealing with machine learning models.

Model Evaluation

To evaluate the predictive performance of the Random Forest model, predictions were generated for the test dataset and compared with the actual survival outcomes.

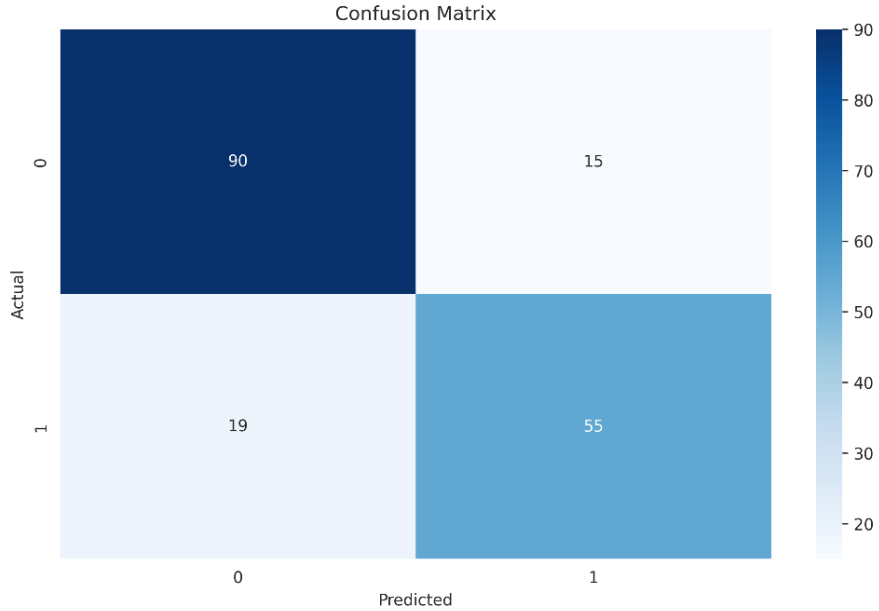

In a confusion matrix we have the number of correct and incorrect predictions of the model. True positives are survivors who are predicted as survivors. True negatives are non-survivors predicted as non-survivors. The model performs well in classification as it accurately predicts most survival cases.

F1 score, precision, and recall were all calculated and presented to show the performance of the model in a better way.

Discussion

According to the results of this tutorial, Random Forest and its model can be interpreted using features importance. The analysis showed that gender and passenger class were the most significant variables for survival.

Consistent with historical accounts of the Titanic disaster, evacuation procedures prioritised women and passengers in higher classes, the results say. The concurrence of model results with historical knowledge highlights the usefulness of machine learning models to elicit meaningful patterns.

Moreover, comparing built-in feature importance to permutation importance shows how different interpretability methods can achieve the same goal in machine learning models. Analysts can achieve a deeper understanding of predictions made by machine learning models with the help of multiple measures of feature importance.

Future Work

Future work could explore additional interpretability techniques such as SHAP values or partial dependence plots to gain deeper insights into how individual features influence predictions. It would also be useful to evaluate the model using larger datasets or alternative ensemble methods such as Gradient Boosting to compare their interpretability and predictive performance.

Limitations

The interpretability tools utilized by Random Forest models are useful; however, there are limitations. The dataset on which model was trained, and model configuration determines feature importance scores. Consequently, if the dataset is changed and if other hyperparameters are used, the importance values may become different.

As well, the fact that some variable is 'important' does not explain how the variable is related to the prediction. To understand complex models better, we may need more sophisticated interpretability methods.

Conclusion

This tutorial demonstrated how Random Forest models can be interpreted using feature importance techniques. Using the Titanic dataset, the analysis identified key variables that influence survival predictions and compared two different approaches for measuring feature importance.

The results highlight the importance of interpretability in machine learning. By examining feature importance scores, analysts can better understand the behaviour of ensemble models and extract meaningful insights from data.

Random Forest models therefore offer a valuable balance between predictive performance and interpretability, making them well suited for many real-world classification problems.

References

Breiman, L. (2001) Random Forests. Machine Learning, 45(1), pp. 5–32.

Géron, A. (2022) Hands-On Machine Learning with Scikit-Learn, Keras and TensorFlow. 3rd ed. O'Reilly Media.

Pedregosa, F. et al. (2011) Scikit-learn: Machine Learning in Python. Journal of Machine Learning Research, 12, pp. 2825–2830.

Fisher, A., Rudin, C. and Dominici, F. (2019) All Models are Wrong, but Many are Useful: Learning a Variable's Importance by Studying an Entire Class of Prediction Models. Journal of Machine Learning Research, 20(177), pp. 1–81.

VanderPlas, J. (2016) Python Data Science Handbook: Essential Tools for Working with Data. O'Reilly Media.

Kaggle (2020) Titanic: Machine Learning from Disaster Dataset. Available at: https://www.kaggle.com/c/titanic